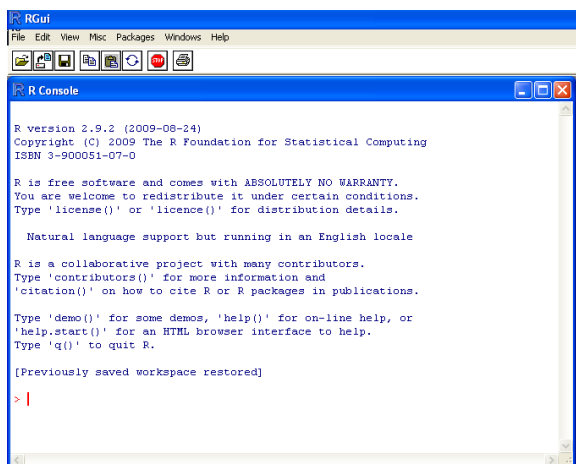


## Getting started with R

This handout gives a very brief introduction to R to give you a rough idea of some basic R features. We assume that R has been installed on your computer.

### Start up R

Start R by double-clicking on the R shortcut on the desktop. A window called the **R Console** will open up.



Type commands after the prompt `>` and then press the **<ENTER>** key.

Note: anything after the pound symbol `#` is a comment—explanatory text that is not executed.

```
> 4*9          # Simple arithmetic
[1] 36
```

If you strike the **<ENTER>** key before typing a complete expression, you will see the *continuation prompt*, the plus sign (+). For example, suppose you wish to calculate  $3 + 2 * (8 - 4)$ , but you accidentally strike the **<ENTER>** key after typing the 8:

```
> 3+2*(8  <ENTER>
+
```

Finish the expression by typing `-4)` after the `+`

```
+ - 4) <ENTER>
[1] 11
```

Create a sequence incrementing by 1:

```
> 20:30
[1] 20 21 22 23 24 25 26 27 28 29 30
```

We will create an object called `dog` and assign it the values 1, 2, 3, 4, 5. The symbol `<-` is the assignment operator.

```
> dog <- 1:5
> dog
[1] 1 2 3 4 5
> dog + 10
[1] 11 12 13 14 15
> 3*dog
[1] 3 6 9 12 15
> sum(dog)          # 1+2+3+4+5
[1] 15
```

The object `dog` is called a *vector*.

If you need to abort a command, press the escape key `<ESC>`. The up arrow key `↑` can be used to recall previous entries.

To obtain help on any of the commands, type the name of the command you wish help on:

```
> ?sum
```

## Importing data

Data for the textbook can be downloaded from the web site

<https://sites.google.com/site/ChiharaHesterberg>

For instance, let's start with the Flight Delays Case Study (see Chapter 2, **Exploratory Data Analysis**, of the text for a description of this data set). For now, we assume the file `FlightDelays.csv` is on the desktop of your computer. We use the `read.csv` command to read it into our R session,

```
> FlightDelays <- read.csv("C:/Users/UserName/Desktop/FlightDelays.csv")
```

or on the Macintosh

```
> FlightDelays <- read.csv("Users/UserName/Desktop/FlightDelays.csv")
```

where `UserName` is the name used to login to this session.

To view the names of the variables in `FlightDelays`:

```
> names(FlightDelays)
```

To view the first part of the data, use the `head` command:

```
> head(FlightDelays)
```

(What do you suppose the `tail` command does?)

The columns are the *variables*. There are two types of variables: *numeric*, for example, `FlightLength` and `Delay` and *factor* (also called *categorical*) (for example `Carrier` and `DepartTime`). The rows are called *observations* or *cases*.

To check the size (dimension) of the data frame, type

```
> dim(FlightDelays)
[1] 4029 10
```

This tells us that there are 4029 observations and 10 columns.

### Tables, bar charts and histograms

The factor variable `Carrier` in the `FlightDelays` data set assigns each flight to one of two *levels*: UA or AA. To obtain the summary of this variable

```
> table(FlightDelays$Carrier)
  AA  UA
2906 1123
```

**Remark:** R is *case-sensitive*! `Carrier` and `carrier` are considered different.

To visualize the distribution of a factor variable, create a *bar chart*:

```
> barplot(table(FlightDelays$Carrier))
```

To compare two categorical variables, the airline (`Carrier`) and whether or not the flight was delayed by more than 30 minutes (`Delayed30`):

```
> table(FlightDelays$Carrier, FlightDelays$Delayed30)

      Delayed30
Carrier   No  Yes
  AA 2513  393
  UA  919  204
```

To see the distribution of a numeric variable, create a histogram.

```
> hist(FlightDelays$Delay)
```

The shape of the distribution of this variable is *right-skewed*.

### Numeric Summaries

Because it is a bit cumbersome to use the syntax `FlightDelays$Delay` each time we want to work with the `Delay` variable, we will create a new object `delay`:

```
> delay <- FlightDelays$Delay
```

```
> mean(delay)
> median(delay)
```

To compute the *trimmed* mean, trimming by 25% on either side,

```
> mean(delay, trim = .25)
```

Other basic statistics:

```
> max(delay)
> min(delay)
> range(delay)
> var(delay)           #variance
> sd(delay)            #standard deviation
> quantile(delay)      #quartiles
```

If you need the population variance (that is, denominator of  $1/n$  instead of  $1/(n-1)$ ),

```
> n <- length(delay)
> (n-1)/n*var(delay)
```

## The `tapply` command

The `tapply` command allows you to compute numeric summaries on values based on levels of a factor variable. For instance, find the mean flight delay length by carrier,

```
> tapply(delay, FlightDelays$Carrier, mean)
```

or the median flight delay length by time of departure

```
> tapply(delay, FlightDelays$DepartTime, median)
```

## Boxplots

Boxplots give a visualization of the 5-number summary of a variable.

```
> summary(delay)           #numeric summary
> boxplot(delay)
```

To compare the distribution of a numeric variable across the levels of a factor variable

```
> tapply(delay, FlightDelays$Day, summary)
> boxplot(Delay ~ Day, data = FlightDelays)
> tapply(delay, FlightDelays$DepartTime, summary)
> boxplot(Delay ~ DepartTime, data = FlightDelays)
```

The `boxplot` command gives the option of using a formula syntax, `Numeric ~ Factor`. In this case, you do not need to use the `$` syntax if you specify the data set.

## Misc. Remarks

- Functions in R are called by typing their name followed by arguments surrounded by *parentheses*: ex. `hist(delay)`. Typing a function name without parentheses will give the code for the function.

```
> sd
```

- We saw earlier that we can assign names to data (we created a vector called `dog`.) Names can be any length, must start with a letter, and may contain letters or numbers:

```
> fish25 <- 10:35
> fish25
```

Certain names are **reserved** so be careful to not use them: `cat`, `c`, `t`, `T`, `F`,...

To be safe, before making an assignment, type the name:

```
> whale
[1] Problem: Object "whale" not found
```

Safe to use `whale`!

```
> whale <- 200
> objects()
> rm(whale)
> objects()
```

- In general, R is space-insensitive.

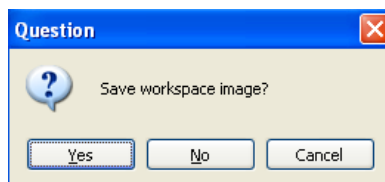
```
> 3 +4
> 3+ 4
> mean(3+5)
> mean ( 3 + 5 )
```

BUT, the assignment operator must not have spaces! `<-` is different from `<` -

- To quit, type

```
> q()
```

You will be given an option to **Save the workspace image**.



Select **Yes**: all objects created in this session are saved to your working directory so that the next time you start up R, if you load this working directory, these objects will still be available. You will not have to re-import `FlightDelays`, for instance, nor recreate `delay`.

You can, for back-up purposes, save data to an external file/disk by using, for instance, the `write.csv` command. See the help file for more information.

---

## Workspace Management

It will be convenient to set up a folder to store the data for this textbook. In the directory where you keep your document, create a folder called **R** and then a subfolder called **MathStats**. The full path depending on the operating system (PC or Mac) will look something like

**C:\Users\UserName\Documents\R\MathStats**,

or

**Users\UserName\Documents\R\MathStats**.

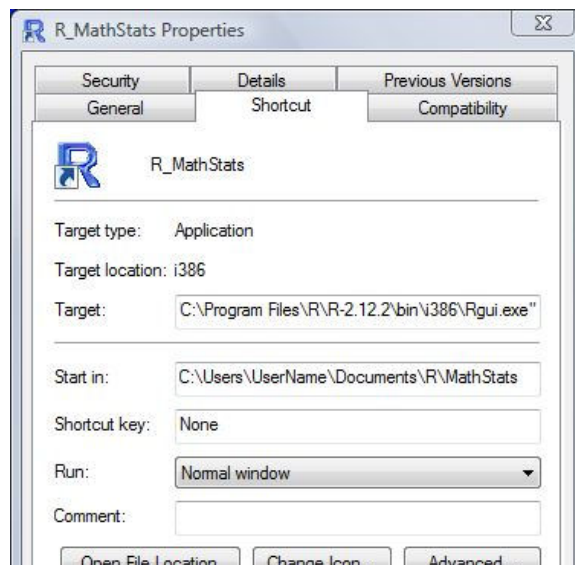
(the exact path will depend on your computer's particular configuration).

Next, we will set up R to start in the above directory.

## Windows

On the desktop of your PC, find the R shortcut, right-click and from the context menu, select **Properties**. In the **Start In:** field, type the path to this folder, for instance

**C:\Users\UserName\Documents\R\MathStats**



and then click **OK**.

Starting R by clicking on this short-cut will automatically load the **MathStats** working directory. To check this,

```
> getwd()
[1] "C:/Users/UserName/Documents/R/MathStats"
```

Copy data files to the working directory. To read them into your R session, you then only need to type the name of the file, rather than a long path. For instance, to read in the data for the Flight Delays Case Study,

```
> FlightDelays <- read.csv("FlightDelays.csv")
```

Rename your shortcut (say to **R.MathStats**) since you may want to set up different working directories for different projects and classes.

## Macintosh

On the Macintosh, dragging the **MathStats** folder icon onto the R icon will start R in the **MathStats** working directory.

Alternatively, if R is already running, then at the menu, under **R > Preferences**, click on **Start-up**. You can change the default working directory to be the **MathStats** folder. Then R will always start with the **MathStats** being the working directory.

Copy data files to the working directory. To read them into your R session, you then only need to type the name of the file, rather than a long path. For instance, to read in the data for the Flight Delays Case Study,

```
> FlightDelays <- read.csv("FlightDelays.csv")
```

## RStudio

**RStudio** is an integrated development environment for R. It is free and can be obtained from <http://www.rstudio.org>. **RStudio** runs on Windows, Macs and Linux. It has a nice interface for managing different projects and workspaces.

---

## Vectors in R

The basic data object in R is the vector. Even scalars are vectors of length 1.

There are several ways to create vectors.

The `:` operator creates sequences incrementing/decrementing by 1.

```
> 1:10
> 5:-3
```

The `seq()` function creates sequences also.

```
> seq(0, 3, by = .2)
> seq(0, 3, length = 15)
> quantile(delay, seq(0, 1, by = .1))      #deciles of delay
```

To create vectors with no particular pattern, use the `c()` function (`c` for `combine`).

```
> c(1, 4, 8, 2, 9)
> x <- c(2, 0, -4)
> x
> c(x, 0:5, x)
```

For vectors of characters,

```
> c("a", "b", "c", "d")
```

or logical values (note that there are no double quotes):

```
> c(T, F, F, T, T, F)
```

The `rep()` command for repeating values:

```
> rep("a", 5)
> rep(c("a", "b"), 5)
> rep(c("a", "b"), c(5, 2))
```

## The “class” attribute

Use `data.class` to determine the class attribute of an object.

```
> state.name
> data.class(state.name)
> state.name == "Idaho"
> data.class(state.name == "Idaho")
> head(FlightDelays$Carrier)
> data.class(FlightDelays$Carrier)
```

## Basic Arithmetic

```
> x <- 1:5
> x - 3
> x*10
> x/10
> x^2
> 2^x
> log(x)

> w <- 6:10
> w
> x*w                      #coordinate-wise multiplication
```



Logical expressions

```
> x < 3
> x == 4
```

## Subsetting

In many cases, we will want only a portion of a data set. For subsetting a vector, the basic syntax is `vector[index]`. In particular, note the use of *brackets* to indicate that we are subsetting.

```
> z <- c(8, 3, 0, 9, 9, 2, 1, 3)
```

The fourth element of `z`:

```
> z[4]
```

The first, third and fourth element,

```
> z[c(1, 3, 4)]
```

All elements *except* the first, third and fourth:

```
> z[-c(1, 3, 4)]
```

To return the values of `z` less than 4, we first introduce the `which` command:

```
> which(z < 4)           # which positions are z values < 4?
> index <- which(z < 4)  # store in index
> z[index]               # return z[c(2, 3, 6, 7)]
```

## The subset command

The `subset` command is used to extract rows of the data that satisfy certain conditions.

If necessary, re-import the `FlightDelays` data set.

Recall that we extracted the variable `delay` using the `$` syntax:

```
> delay <- FlightDelays$Delay
```

The `subset` command can also be used:

```
> delay <- subset(FlightDelays, select = Delay, drop = TRUE)
```

The `select =` argument indicates which column to choose.

The `drop = TRUE` argument is need to create a `vector`. Compare the output of

```
> subset(FlightDelays, select = Delay, drop = TRUE)
```

to

```
> subset(FlightDelays, select = Delay)
```

The second output (omitting the `drop = TRUE` argument) is a `data.frame` object (more on data frames later).

Suppose you wish to extract the flight delay lengths for all days except Monday and then find the mean delay length:

```
> delay2 <- subset(FlightDelays, select = Delay, subset = Day != "Mon",
  drop = TRUE)
> mean(delay2)
```

The `subset =` argument is a logical expression indicating which rows to keep. We want those days not equal to Monday.

To extract the delay lengths for Saturday and Sunday only

```
> delay3 <- subset(FlightDelays, select = Delay,
  subset = (Day == "Sat" | Day == "Sun"), drop = TRUE)
> mean(delay3)
```

The vertical bar `|` stands for “or.”

Suppose you want to find those observations when the delay length was greater than the mean delay length. We’ll store this in a vector called `index`.

```
> index <- which(delay > mean(delay))
> head(index)
[1]  2 10 12 14 15 16
```

Thus, observations in rows 2, 10, 12, 14, 15, 16 are the first six that correspond to flights that had delays that were larger than the average delay length.

### Vectorized operators

<code>==</code>	equal to
<code>!=</code>	not equal to
<code>&gt;</code>	greater than
<code>&gt;=</code>	greater than or equal to
<code>&lt;</code>	less than
<code>&lt;=</code>	less than or equal to
<code>&amp;</code>	vectorized AND
<code> </code>	vectorized OR
<code>!</code>	not

**Programming Note:** The vectorized AND and OR are for use with vectors (when you are extracting subsets of vectors). For control in programming (ex. when writing `for` or `if` statements), the operators are `&&` and `||`.

### Misc. commands on vectors

```
> length(z)      # number of elements in z
> sum(z)         # add elements in z
> sort(z)        # sort in increasing order
```

The `sample` command is used to obtain random samples from a set.

For instance, to permute the numbers  $1, 2, \dots, 10$ :

```
> sample(10)
```

To obtain a random sample from  $1, 2, \dots, 10$  of size 4 (without replacement):

```
> sample(10, 4)
```

Without replacement is the default; if you want to sample with replacement:

```
> sample(10, 4, replace = TRUE)
```

The built-in vector `state.name` contains the 50 states.

```
> state.name
```

Draw a random sample of 20 states:

```
> sample(state.name, 20)
```

If you want to sample with replacement,

```
> sample(state.name, 20, replace=TRUE)
```

Suppose you wish to create two vectors, one containing a random sample of 20 of the states, the other vector containing the remaining 30 states.

We obtain a random sample of size 20 from  $1, 2, \dots, 50$  and store in the vector `index`:

```
> index <- sample(50, 20)
```

See what this sample looks like:

```
> index
```

`tempA` will contain the random sample of 20 states:

```
> tempA <- state.name[index]
```

`tempB` will contain the remaining 30 states. Recall that in subsetting, the negative of a number means to exclude that observation.

```
> tempB <- state.name[-index]
```

```
> tempA
```

```
> tempB
```

## Data Frames in R

Most data will be stored in data frames, rectangular arrays which usually are formed by combining columns of vectors. `FlightDelays` is an example of a data frame.

```
> data.class(FlightDelays)
```

### Subsetting data frames

For subsetting a data frame, use the syntax  
`data[row.index, column.index]`.

For instance, row 5, column 3:

```
> FlightDelays[5, 3]
```

or rows 1 through 10, columns 1 and 3:

```
> FlightDelays[1:10, c(1, 3)]
```

or all rows *except* 1 through 10, and keep columns 1 and 3:

```
> FlightDelays[-(1:10), c(1, 3)]
```

To extract *all* rows, but columns 1 and 3

```
> FlightDelays[, c(1, 3)]
```

and rows 1:100 and *all* columns

```
> FlightDelays[1:100, ]
```

Use the `subset` command to extract rows based on some logical condition.

Create a subset of just the Tuesday data:

```
> DelaysTue <- subset(FlightDelays, subset = Day == "Tue")  
> head(DelaysTue)
```

Create a subset of just the Tuesday data and columns 1,6,7:

```
> DelaysTue <- subset(FlightDelays, Day == "Tue", select = c(1, 6, 7))  
> head(DelaysTue)
```

## Plots in R

There are three basic plotting functions in R: high-level plots, low-level plots, and the layout command `par`. Basically, a high-level plot function creates a complete plot and a low-level plot function adds to an existing plot, that is, one created by a high-level plot command.

### High-Level Plot Functions

Some of the basic plot functions include:

Function	Description
<code>plot</code>	scatter/line plot
<code>hist</code>	histogram
<code>barplot</code>	barplot
<code>boxplot</code>	boxplot
<code>qqnorm</code>	normal-quantile

Download the example data set **States03** from <http://apps.carleton.edu/curricular/math/resources/rcomputing/>, then import it into your session.

If you are using RStudio, then at the menu, **Tools > Import Dataset > From Text File...** and navigate to the location of the file.

Or at the command line, use `read.csv` to import the data:

```
> States03 <- read.csv("States03.csv")
```

(exact path will vary depending on where you saved the file).

Basic single variable plots:

```
> barplot(table(States03$Region))
> hist(States03$Poverty)
```

To create a scatter plot, there are two approaches:

```
> plot(States03$Unemp, States03$Poverty, xlab = "Unemployment", ylab = "Poverty")
> plot(Poverty ~ Unemp, data = States03, xlab = "Unemployment", ylab = "Poverty")
```

In the first approach, provide the `plot` command with the x-variable, then the y-variable. In the second approach, if the data are contained in a data frame, then provide the names of the variable  $Y \sim X$  along with the name of the data frame.

High-level functions may also take optional arguments that enhance the plot.

```
> hist(States03$Poverty, main = "Poverty", xlab = "percent",
      xlim = c(0, 24), ylim = c(0, 20))
> plot(1:19, 1:19, pch = 1:19, xlab = "x", ylab = "y")
> pie(rep(1, 8), col = 1:8)
```

Option	Description
pch	point character (pch=1, 2, ...)
lty	line type (lty=1, 2, ...)
lwd	line thickness (lwd= 1, 2,...)
col	color (col="red", "blue",...)
xlim	x-axis limits: xlim=c(min,max)
ylim	y-axis limits
xlab	x-axis label: xlab="my label"
ylab	y-axis label
main	main title
sub	sub title

To plot smooth curves, use the `curve` command. The first argument must be an expression in terms of  $x$ :

```
> curve(x^2, from = 0, to = 2)
> curve(cos(x), from = 0, to = pi)
> curve(cos(x), from = 0, to = pi, lty = 4, col = "red")
```

## Low-level Plot Functions

Low-level plot functions can be executed only after a high-level plot has been created. For example,

```
> plot(Poverty ~ Unemp, data = States03, xlab = "Unemployment", ylab = "Poverty")
> abline(v = mean(States03$Unemp), lty = 2) #vertical line at mean unemployment rate,
> text(30, 18, "mean unemployment rate") #text at (30, 18)
> title("Data from 2003")
```

The `abline` function has several options:

`abline(3, 5)` adds the straight line  $y = 3 + 5x$

`abline(v = 2)` adds the vertical line,  $x = 2$

`abline(h = 0)` adds the horizontal line,  $y = 0$

```
> plot(Poverty ~ ColGrad, data = States03, col = "blue", pch = 19, xlab = "College grad (%)",
      ylab = "Poverty (%)")
> points(Uninsured ~ ColGrad, data = States03, col = "red", pch = 19)
> mtext("Percent uninsured", side = 4)
> legend("bottomleft", legend = c("Y: Poverty", "Y: Uninsured"), col = c("blue", "red"),
      pch = c(16, 16))
```

You can also use different plotting symbols for different levels of a factor variable:

## Probability Distributions in R

### Continuous

Distributions	root
beta	<b>beta</b>
Cauchy	<b>cauchy</b>
chi-square	<b>chisq</b>
exponential	<b>exp</b>
F	<b>f</b>
gamma	<b>gamma</b>
normal	<b>norm</b>
student's t	<b>t</b>
uniform	<b>unif</b>
Weibull	<b>weibull</b>

In the continuous case, **droot** returns the density, **proot** a cumulative probability, **qroot** a quantile, **rroot** a random number.

### Probability

If  $X$  follows  $N(0, 1)$ , then to find  $P(X \leq 1.25) = \Phi(1.25)$ , that is, the amount of area under the standard normal density curve to the left of  $x = 1.25$ ,

```
> pnorm(1.25)
```

By default, the **norm** function assumes  $\mu = 0, \sigma = 1$  (that is, you are working with the standard normal distribution). For other means and standard deviations, specify them in the argument. For example, if  $X \sim N(\mu = 2, \sigma = 3)$ , then to find  $F(2.8) = P(X \leq 2.8)$ ,

```
> pnorm(2.8, 2, 3)
```

If  $X$  follows a chi-square distribution with 25 degrees of freedom then to compute  $F(13.9) = P(X \leq 13.9)$ ,

```
> pchisq(13.9, 25)
```

If  $X$  follows an exponential distribution with parameter  $\lambda = 10$ , then to compute  $P(X > 4)$ ,

```
> 1 - pexp(4, 10)
```

or

```
> pexp(4, 10, lower.tail=FALSE)
```

If  $T$  follows a t-distribution with 7 degrees of freedom, then to find the probability that  $T \leq 3.9$ , type

```
> pt(3.9, 7) # pt(t-value, d.f)
```

### Quantiles

To find the 25<sup>th</sup> percentile, that is, the value  $q$  such that  $P(X \leq q) = .25$  for  $X$  from  $N(0, 1)$ ,

```
> qnorm(.25)
[1] -0.6744898
```

In other words, the amount of area under the pdf to the left of  $x = -0.6744$  is 0.25. Or, if  $F$  denotes the cdf of the distribution, then  $F^{-1}(0.25) = -0.6744$ .

The .75 quantile for  $N(2, 3)$  can be found by

```
> qnorm(.75, 2, 3)
[1] 4.023469
```

In other words, the amount of area under the density curve and to the left of  $x = 4.023469$  is .75, or if  $F$  denotes the cdf, then  $F^{-1}(0.75) = 4.023469$ .

For  $T$  from a t-distribution with 13 degrees of freedom, to find value  $t$  such that  $P(T > t) = .025$ , which is equivalent to  $F(t) = P(T \leq t) = 0.975$ , type

```
> qt(.975, 13)
```

### Random numbers

To generate 100 random numbers from the normal distribution  $N(0, 1)$ , type

```
> rnorm(100)
> x <- rnorm(100)
> hist(x)
```

Ten random numbers from the chi-square distribution with 23 degrees of freedom,

```
> rchisq(10, 23)
```

### Plotting the density curve (pdf)

To plot the pdf for  $N(0, 1)$  for  $-3 \leq x \leq 3$ , use the `curve` function with the pdf `dnorm` provided as an argument.

```
> curve(dnorm(x), from = -3, to = 3)
> w <- rnorm(50) # random sample from N(0,1)
> hist(w, freq = TRUE) # scale to area 1
> curve(dnorm(x), add = TRUE) # impose normal density
```



```
> hist(w, freq = TRUE, ylim = c(0, .5)) # widen y-axis range
> curve(dnorm(x), add=TRUE)
```

To plot the pdf for the chi-square distribution with 14 degrees of freedom,

```
> curve(dchisq(x, 14), from = 0, to = 20)
```

## Discrete

Distribution	root
binomial	<b>binom</b>
geometric	<b>geom</b>
hypergeometric	<b>hyper</b>
negative binomial	<b>nbinom</b>
Poisson	<b>pois</b>

Preface each of the above roots with either **d**, **p**, **q** or **r**.

**droot** is the probability mass function so returns a probability, **proot** returns a cumulative probability (cmf), and **qroot** returns a quantile, and **rroot** returns a random number.

The quantile function is the inverse of the CDF,  $F(t) = P(X \leq t) = \sum_{k \leq t} P(X = k)$ .

### Example Binomial

Suppose you have a biased coin that has a probability of 0.8 of coming up heads.

The probability of getting 5 heads in 16 tosses of this coin is

```
> dbinom(5,16, .8)
```

Check this answer by calculating directly

$$\binom{16}{5} .8^5 \cdot .2^{11},$$

```
> choose(16,5)*.8^5*.2^11
```

The probability of getting at most 5 heads in 16 tosses is

```
> pbinom(5,16, .8)
```

In other words, `pbinom(5, 16, .8)` is computing: `dbinom(0,16,.8)+dbinom(1,16,.8)+dbinom(2,16,.8)+dbinom(3,16,.8)+dbinom(4,16,.8)+dbinom(5,16,.8)`

The 0.25 quantile is

```
> qbinom(.25,16,.8)
```

```
[1] 12
```

This is the smallest number of successes such that the probability of at most this many successes is greater than or equal to .25.

Check this:

```
> pbinom(11, 16, .8)
> pbinom(12, 16, .8)
```

### Example (cont.) Geometric

Find the probability of getting the first head on the fourth toss. This is the geometric distribution. The arguments to `geom` are `geom(failures, p)`.

```
> dgeom(3, .8)
```

The probability that the first head occurs on one of the first four tosses (that is, on the first, second, third or fourth toss) is

```
> pgeom(3, .8)
```

### Example Poisson

Suppose a certain region of California experiences about 5 earthquakes a year. Assume occurrences follow a Poisson distribution. What is the probability of 3 earthquakes in a given year?

Here  $\lambda = 5$

```
> dpois(3,5)
```

Check the answer:

```
> 5^3*exp(-5)/(3*2)
```

### Random numbers

To generate random numbers from a particular distribution, preface the root name with an `r`.

For example, we continue our previous example of a biased coin with  $p = .8$  of coming up heads. Toss this coin 25 times. The command `rbinom(1,25,.8)` will return a random number of successes.

```
> rbinom(1,25,.8)
```

Now, let's run this experiment 10 times (that is, we do 10 sets of tossing a coin 25 times) and record the number of successes.

```
> set.seed(0)
```

This sets the seed for the random number generator so that we all get the same results.

```
> heads <- rbinom(10, 25, .8)
> heads
```

```
[1] 17 19 21 18 20 18 22 18 22 17 #output will vary
```

In the first experiment of tossing the coin 25 times, 17 heads occurred. In the second experiment of tossing the coin 25 times, 19 heads occurred, etc.

```
> table(heads)
```

```
> barplot(table(heads))
```

Repeat the above, except now run the experiment 100 times.

```
> heads2 <- rbinom(100, 25, .8)
```