

1. The data frame `HSWRESTLER`, (from package `PASWR2`) contains information on nine variables for a group of 78 high school wrestlers that was collected by the human performance lab at Appalachian State University. It is of interest to predict wrestler's hydrostatic fat (`hwfat`) using predictors `age`, `ht`, `wt`, `abs`, `triceps` and `subscap`. Split the data set into a training set and a test set (50%). Use `set.seed(1)` each time you need to use `sample()` or `cv.glmnet()` functions.
 - (a) Fit a linear model using least squares on the training set, and report the test error obtained.
 - (b) Fit a ridge regression model on the training set, with λ chosen by cross-validation (use 15-fold cross validation). Report the test mspe.
 - (c) Fit a lasso model on the training set, with λ chosen by cross-validation (use 15-fold CV). Report the test error obtained.

```
# hw3q1.r

library(PASWR2)
library(glmnet)

# a) training and test sets

d0 = HSWRESTLER[,1:7]
n = nrow(d0)
p = ncol(d0)-1

set.seed(1)
train=sample(1:n, n/2)
dtrain = d0[train,]
dtest = d0[-train,]

# (b) linear model
#=====

m1 = lm( hwfat ~ ., dtrain )
Y_hat = predict(m1,dtest )

mspe = mean((dtest$hwfat - Y_hat)^2)
mspe
# 14.0963
```

```

# (c) ridge regression
#=====

Y = dtrain$hwfat
MM = model.matrix(hwfat ~ ., data=dtrain)[,-1]

newx=model.matrix(hwfat~.,dtest)[,-1]

set.seed(1)
cv.out = cv.glmnet(MM,Y, alpha=0,nfolds=15)
bestlam = cv.out$lambda.min
bestlam
# 0.994199

ridge.mod = glmnet(MM,Y,alpha=0)
Y_hat = predict(ridge.mod,s=bestlam,newx)

mspe2 =mean((dtest$hwfat-Y_hat)^2)
mspe2
# 14.5059

# refit with full data set

Y = d0$hwfat
MM = model.matrix(hwfat~.,data=d0)[,-1]

out=glmnet(MM,Y,alpha=0)
options(scipen=999)      # disable scientific notation
options(digits=4)
ridge.coef=predict(out,type="coefficients",s=bestlam) [1:p,]
ridge.coef
# (Intercept)      age      ht      wt      abs      triceps      subscap      ta
# 15.408559   -0.430032  -0.128271  0.006445  0.136425  0.232287  0.128201  0.19

```

```

# (d) lasso regression
#=====

Y = dtrain$hwhfat
MM = model.matrix(hwhfat ~ ., data=dtrain)[,-1]

set.seed(1)
cv.out = cv.glmnet(MM,Y,alpha=1,nfolds=15)
bestlam = cv.out$lambda.min
bestlam
# 0.04948

lasso.mod = glmnet(MM,Y,alpha=1)
Y_hat = predict(lasso.mod,s=bestlam,newx)

mspe3 =mean((dtest$hwhfat-Y_hat)^2)
mspe3
# 13.77

# refit with full data set

Y = d0$hwhfat
MM = model.matrix(hwhfat~.,data=d0)[,-1]

out=glmnet(MM,Y,alpha=1)
options(digits=6)
lasso.coef=predict(out,type="coefficients",s=bestlam) [1:p,]
lasso.coef
#(Intercept)      age      ht      wt      abs      triceps
# 15.214113  -0.382793  -0.104636  0.000000  0.341984  0.397083

lasso.coef[lasso.coef!=0]
#(Intercept)      age      ht      abs      triceps
# 15.214113  -0.382793  -0.104636  0.341984  0.397083

```

2. Generate simulated data, and will then use this data to perform best subset selection. Generate values of a predictor X of length $n = 100$, using $X = \text{rnorm}(n)$. Generate values of a noise vector ϵ of length $n = 100$ using $\epsilon = 0.1 * \text{rnorm}(n)$. Generate a response vector Y of length $n = 100$ using

$$Y = 1 - 0.1X + 0.05X^2 + 0.75X^3 + \epsilon$$

- (a) Use `regsubsets()` to choose the best model containing the predictors X, X^2, \dots, X^{10} . What is the best model obtained according to adjusted R^2 ?
- (b) Fit a lasso model to the simulated data, again using X, X^2, \dots, X^{10} as predictors. Use 10-fold cross-validation to select the optimal value of λ . Report the test error.

```
set.seed(1)
n = 100
x = rnorm(n)
epsilon = 0.1 * rnorm(n)

beta_0 = 1.0
beta_1 = -0.1
beta_2 = +0.05
beta_3 = 0.75
y = beta_0+beta_1*x+beta_2*x^2+beta_3*x^3+epsilon
d0 = data.frame(y,x,x2=x^2,x3=x^3,x4=x^4,x5=x^5,x6=x^6,x7=x^7,x8=x^8,x9=x^9,x10=x^10)

# split training/test sets

set.seed(1)
train=sample(1:n, n/2)
test = (-train)
dtrain = d0[train,]
dtest = d0[test,]

# a) regsubsets
library(leaps)

models = regsubsets(y~.,d0,nvmax=11)
summary(models)
# Selection Algorithm: exhaustive
#      x  x2 x3 x4 x5 x6 x7 x8 x9 x10
# 1 ( 1 ) " " " " "*" " " " " " " " " " "
# 2 ( 1 ) " " " " "*" " " "*" " " " " " " " "
# 3 ( 1 ) "*" " " "*" "*" " " " " " " " " " "
# 4 ( 1 ) "*" "*" "*" " " "*" " " " " " " " "
# 5 ( 1 ) "*" "*" "*" " " " "*" "*" " " " " " "
# 6 ( 1 ) "*" "*" "*" " " " " " " "*" "*" "*" " "
# 7 ( 1 ) "*" " " "*" "*" " " " "*" "*" "*" " " "*"
# 8 ( 1 ) "*" "*" "*" "*" " " " "*" " " "*" "*" "*"
# 9 ( 1 ) "*" "*" "*" "*" "*" "*" " " " "*" "*" "*"
# 10 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"

summary(models)$adjr2
```

```

# 0.996226 0.997172 0.997571 0.997600 0.997586 0.997565 0.997556 0.997534 0.997507 0.997484
a=summary(models)$adjr2
which.max(a)    # 4
# best model with predictors x,x2,x3,x5

#b) lasso
#=====
library(glmnet)

MM = model.matrix(y~.,dtrain)
ytrain = dtrain$y

# best lambda from train set
set.seed(1)
cv.out = cv.glmnet(MM,ytrain,alpha=1)
bestlam = cv.out$lambda.min
bestlam
# 0.0255913

lasso.mod = glmnet(MM,ytrain,alpha=1)

# predict test set
newx = model.matrix(y~.,dtest) #[, -1]
yhat = predict(lasso.mod,s=bestlam,newx)    # error

ytest = dtest$y
mspe = mean((ytest-yhat)^2)
mspe
# 0.0103474

# refit
ytrain = d0$y
MMfull = model.matrix(y~.,d0)
out=glmnet(MMfull,d0$y,alpha=1)

lasso.coef=predict(out,type="coefficients",s=bestlam)[1:12,]
lasso.coef
lasso.coef[lasso.coef!=0]
# (Intercept)          x3          x5
# 1.03708472 0.64348867 0.01368301

```

3. A real estate appraiser is interested in predicting residential home prices in a mid-western city as a function of various features. For that purpose a regression model is to be constructed from a sample of 522 houses. Use the `homes.xls` data set from blackboard. Consider the predictors x_1 : lot size (square feet), x_2 : area (square feet), x_3 : number of bedrooms, x_4 : number of bathrooms, x_5 : year of construction, x_6 : garage size (number of cars). Split the data set into a training set and a test set (50%).
- Fit a ridge regression model on the training set, with λ chosen by cross-validation (use 15-fold cross validation). Report the test mspe.
 - Fit a lasso model on the training set, with λ chosen by cross-validation (use 15-fold CV). Report the test error obtained.
 - Predict the price when all predictors are equal to their median values using both the ridge regression and the lasso models.

```
# hw3q3.r

library(glmnet)
setwd("C:/Users/USC Guest/Downloads2")
d0=read.csv("homes.csv",header=T)
d1 = subset(d0,select=c(price,lotsize,area,beds,baths,year,garage))

# a) training and test sets

set.seed(1)
n = nrow(d1)
p = ncol(d1)-1

train=sample(1:n, n/2)
test = (-train)
dtrain = d1[train,]
dtest = d1[test,]

# linear model (not required)
#=====

m1 = lm( price ~ ., dtrain )
Y_hat = predict(m1,dtest )

mspe = mean((dtest$price - Y_hat)^2)
mspe
# 5048263624
```

```

# (a) ridge regression
#=====

Y = dtrain$price
newx=model.matrix(price~.,dtest)
newx=model.matrix(price~.,dtest)[,-1]
MM = model.matrix(price ~ ., data=dtrain)[,-1]

head(dtrain)
head(MM)

set.seed(1)
cv.out = cv.glmnet(MM,Y, alpha=0,nfolds=15)
bestlam = cv.out$lambda.min
bestlam
# 12660.1

ridge.mod = glmnet(MM,Y,alpha=0)

# model from training set
ridge.coef=predict(ridge.mod,type="coefficients",s=bestlam) [1:p,]
ridge.coef
#      (Intercept)      lotsize          area          beds          baths          year
# -4157738.62649      1.25728      96.29050     -3482.30449      9980.27572      2096.76099

# Prediction at the median with training data set model
#-----
apply(d1,2,median)
# price lotsize   area   beds   baths   year garage
# 229900  22200   2061     3     3   1966     2

# order matters
newval2= data.frame(lotsize=22200,area=2061,beds=3,baths=3,year=1966,garage=2)
newval2 = as.matrix(newval2)
newval2
#      lotsize area beds baths year garage
# [1,]  22200 2061    3    3 1966     2
Y_hat = predict(ridge.mod,s=bestlam,newval2)
Y_hat
#           1
# [1,] 256792

# mspe
Y_hat = predict(ridge.mod,s=bestlam,newx)
mspe2 =mean((dtest$price-Y_hat)^2)
mspe2

# 5280252885

# refit

```

```
Y = d1$price
MM = model.matrix(price~.,data=d1)[,-1]

out=glmnet(MM,Y,alpha=0)
options(scipen=999)      # disable scientific notation
options(digits=4)
ridge.coef=predict(out,type="coefficients",s=bestlam) [1:p,]
ridge.coef
# (Intercept)      lotsize        area        beds        baths        year
#-3290230.723        1.428        104.289       -8129.583       14828.204       1642.448

# Prediction at the median with refitted model
#-----
Y_hat = predict(out,s=bestlam,newval2)
Y_hat
#           1
#[1,] 258975
```



```

# (b) lasso regression
#=====

Y = dtrain$price
MM = model.matrix(price~.,data=dtrain)[,-1]

set.seed(1)
cv.out = cv.glmnet(MM,Y,alpha=1,nfolds=15)
bestlam = cv.out$lambda.min
bestlam    # 574.1

lasso.mod = glmnet(MM,Y,alpha=1)

# Prediction at the median with training data set model
Y_hat = predict(lasso.mod,s=bestlam,newval2)
Y_hat
#           1
# [1,] 251434

# mspe
Y_hat = predict(lasso.mod,s=bestlam,newx)
mspe3 =mean((dtest$price-Y_hat)^2)
mspe3
# 5084835702

# refit with full data set

Y = d1$price
MM = model.matrix(price~.,data=d1)[,-1]

out=glmnet(MM,Y,alpha=1)
options(digits=6)
lasso.coef=predict(out,type="coefficients",s=bestlam) [1:p,]
lasso.coef

#      (Intercept)      lotsize          area          beds          baths          year
# -3542559.21938      1.52412      125.29191     -12118.07170      7565.80255      1766.82300

# No coefficients equal to zero

# Prediction at the median with refitted model
#-----
Y_hat = predict(out,s=bestlam,newval2)
Y_hat
#           1
# [1,] 253919

```